

CSE 536: Advanced Operating Systems

Assignment 0: Getting Started (ungraded)

1 Brief Description of the xv6 OS and RISC-V

xv6 is an education-focused operating system (OS) designed by folks from MIT, namely Franz Kaashoek, Robert Morris, and Russ Cox (alongside contributions from many other individuals). Like Linux and MacOS, xv6 follows the structure of UNIX-based operating systems, and it is a loose reimplementation of Dennis Ritchie and Ken Thompson's UNIX version 6 (v6). UNIX-based OSs are commonplace today, and each of them typically consists of 3 main components: the kernel, the shell, and the applications/file system. Kernel manages the system operations (e.g., memory management, resource allocation, etc.) and the Shell is used to interact with the kernel.




Figure 1: UNIX system architecture ([source](#)).

Suggested reading

[The UNIX Time-Sharing System](#), Dennis M. Ritchie and Ken L. Thompson, Bell System Technical Journal 57, number 6, part 2 (July-August 1978) pages 1905-1930.

[xv6: A Simple, Unix-like Teaching Operating System](#), Russ Cox, Franz Kaashoek, and Robert Morris

xv6 is designed as a *monolithic* kernel, meaning all functionality of the OS is merged into one large program (please check notes from the 2nd lecture for more details). While modern monolithic kernels are terribly large (e.g., Linux contains more than 31 million code lines), the same cannot be said for xv6, since it is a fairly simple and self-contained kernel (approximately 6 thousand code lines).

Question.1

What are the advantages and disadvantages of monolithic kernels?

In this class, the implementation of xv6 we are concerned with is written in C and designed for the RISC-V architecture. This architecture was introduced at UC Berkeley in the early 2000s. RISC stands for Reduced Instruction Set Computing which at a high-level means that the architecture does not provide CPU instructions that perform complex tasks (e.g., run a cryptographic algorithm like AES). A famous example of RISC architecture is ARM, which is widely used in mobile phones (e.g., Qualcomm Snapdragon) and

in modern Apple machines. In contrast to RISC, there is a CISC or Complex Instruction Set Computing architecture, which is designed to allow the CPU to perform complex tasks. A famous example of CISC architecture is the widely-used x86 architecture co-designed by Intel and AMD.

Tidbit

You might ask the following question: “*why are we using xv6 instead of a well-known OS?*” For better or worse, your instructor’s belief is that understanding core OS concepts is too complex on modern OSs like Linux (and heavens forbid Windows or MacOS). The goal of this course is to help you understand core concepts, instead of having you understand the vagaries of modern OSs.

2 Running xv6 on ASU Apporto Instances (or Personal Machines)

The xv6 operating system is designed for the RISC-V architecture (as mentioned in the previous paragraph). Hence, generally to run the xv6 OS, we require (1) a CPU capable of running RISC-V instructions and (2) a compiler toolchain capable of producing RISC-V assembly code. Requirement (2) is straightforward (we can easily download and set up a RISC-V compiler), but requirement (1) is very tricky since none of our CPUs support RISC-V instructions natively (more than likely you are running an x86 or ARM CPU). Hence, to fulfill the requirement (1), we will emulate a RISC-V CPU using a system emulator called QEMU. It emulates the machine’s CPU processor through dynamic binary translation and provides a set of different hardware and device models for the machine, enabling it to run a variety of guest operating systems. This is perfect for our use case to learn RISC-V architecture and implement OS functionality.

The following explains how to set up xv6 on your personal machines and Apporto cloud instances.

2.1 Setting up xv6 on your Linux, WSL, or Mac machines

Please follow the steps below to install on a Linux/WSL machine. For a Mac machine, please follow steps in the repository’s README.

1. Clone the CSE 536 GitHub repository and checkout the getting-started branch.

```
git clone https://github.com/ASTERISC-ASU/cse536-release.git  
git checkout getting-started
```
2. Download and install the RISC-V GNU compiler toolchain.

```
cd install/linux-wsl && ./linux-toolchain.sh
```
3. Download and install the RISC-V capable QEMU.

```
cd install/linux-wsl && ./linux-qemu.sh
```
4. Export the installed components to PATH so that you can access them from shell.

```
cd install/linux-wsl && source .add-linux-paths
```

Tidbit

Step 4 only exports the toolchains and QEMU to PATH for the currently running shell. If you exit this shell, you will have to repeat step 4. You can avoid this repetition (Google it ;)).

5. Clone and run xv6.

```
git clone https://github.com/mit-pdos/xv6-riscv.git
```

```
cd xv6-riscv
make qemu
```

6. You should see a screen like the following, which basically opens the xv6 shell.

```
> make qemu
qemu-system-riscv64 -machine virt -bios none -kernel kernel/kernel -m 128M -smp
 3 -nographic -global virtio-mmio.force-legacy=false -drive file=fs.img,if=none
,format=raw,id=x0 -device virtio-blk-device,drive=x0,bus=virtio-mmio-bus.0

xv6 kernel is booting

hart 2 starting
hart 1 starting
init: starting sh
$ █
```

Figure 2: Terminal view - make qemu

The `make qemu` command creates a QEMU virtual machine (VM) and runs the xv6 kernel inside the VM. There are different parameters provided, which you can inspect by looking under the line `(make qemu)` in the Makefile. More details in the next section.

Once booted, you're in the xv6 kernel which shows you a shell terminal (look at '\$...'). Play around and try commands like `ls`, `cat`, etc. They should work similarly to a normal shell. To find out all the commands supported by the shell, run '`ls`'.

Tidbit

To exit xv6 and go back to your host terminal shell, execute `Ctrl+A` then `X`.

Please use the following [link](#) as additional reference. If you run into issues, reach out to the TAs for help.

2.2 Opening Apporto on Canvas

Apporto is a virtual cloud environment with many different OS instances (e.g., Linux, Windows, MacOS) provided by ASU to students through canvas.

1. Go to your ASU canvas website and select the canvas for CSE 536.
2. On your canvas, there should be an "Apporto Virtual Lab" on the left tab.
3. Apporto should open an Ubuntu-based operating system environment for you.

Important note

Apporto instances are stateless – all stored information is deleted once you log out of the instance. You also get logged out automatically after a short period of inactivity (e.g., 5-10 minutes). Hence, it is VERY IMPORTANT to save your assignments/code in a network drive BEFORE you log out of the instance. In fact, it is a good practice to save periodically (e.g., after you are done with one part

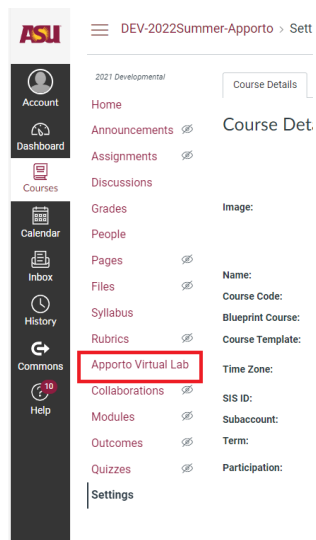


Figure 3: Finding Apporto on Canvas.




Figure 4: An Apporto instance.

of the assignment).

2.3 Running xv6 on Apporto

In the Apporto cloud instances provided to you, QEMU and RISC-V GNU toolchains should be pre-installed and configured to run xv6 beforehand.

Important note

As of Jan 12, 2023, Apporto is not correctly configured to run xv6. In particular, it does not currently have the packages needed to install QEMU or the RISC-V GNU toolchain, and we're not allowed root privileges. This is likely to change in the coming days; please stand-by. Once Apporto is correctly configured, the steps below will become useful.

Tidbit

To test out if QEMU and RISC-V toolchains are correctly installed on the Apporto instance, please run the following commands: `qemu-system-riscv` and `riscv64-unknown-elf-gcc`.

1. Open a terminal within the Ubuntu instance.
2. Clone the xv6 RISC-V OS repository from GitHub.
`git clone https://github.com/mit-pdos/xv6-riscv.git`
3. Navigate into the cloned directory.
`cd xv6-riscv`
4. Run the xv6 OS.
`make qemu`

2.4 Using the Network Storage on Apporto

To be explained in assignment #1.

3 Debugging xv6 using the GNU debugger (GDB)

A debugger is your friend when dealing with systems-level programming (aka designing an OS), and GDB is your best friend. To run xv6 in debugging mode with GDB, please follow the steps below:

1. Instead of running `make qemu` on step 4 in [subsection 2.3](#), you can run `gdb` for xv6 using `make qemu-gdb`. This will produce an output similar to the one below.

```
> make qemu-gdb
*** Now run 'gdb' in another window.
qemu-system-riscv64 -machine virt -bios none -kernel kernel/kernel -m 128M -smp
3 -nographic -global virtio-mmio.force-legacy=false -drive file=fs.img,if=none
,format=raw,id=x0 -device virtio-blk-device,drive=x0,bus=virtio-mmio-bus.0 -S -
gdb tcp::26000
█
```

Figure 5: Terminal view - make qemu-gdb

Essentially, the system starts in debug mode and waits for a GDB instance to be attached.

2. On a separate terminal window, navigate to the `xv6-riscv` folder.
3. Then, run `riscv64-unknown-elf-gdb` to attach a GDB instance to the QEMU VM.

```

> riscv64-unknown-elf-gdb
GNU gdb (GDB) 12.1
Copyright (C) 2022 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law.
Type "show copying" and "show warranty" for details.
This GDB was configured as "--host=x86_64-pc-linux-gnu --target=riscv64-unknown-elf".
Type "show configuration" for configuration details.
For bug reporting instructions, please see:
<https://www.gnu.org/software/gdb/bugs/>.
Find the GDB manual and other documentation resources online at:
<http://www.gnu.org/software/gdb/documentation/>.

For help, type "help".
Type "apropos word" to search for commands related to "word".
The target architecture is set to "riscv:rv64".
warning: No executable has been specified and target does not support
determining executable automatically. Try using the "file" command.
0x0000000000000100 in ?? ()
(gdb) █

```

Figure 6: Terminal view - GDB debugger

Tidbit

It is important to run GDB in the xv6-riscv folder because we have set up various configurations for the debugger in the .gdbinit script in the folder. If you're interested, check out that script.

Now you can debug execution by using the following commands on the gdb terminal:

- `continue` (or `c`) to resume execution until breakpoint.
- `b func` to set a breakpoint at a certain function (func).
- `b filename:line` to set a breakpoint at a certain line of a file.
- `si` to single-step the execution by running one instruction and stopping.
- `s` to single-step the execution by running one C code line and stopping.
- `info registers` will show you the value of various CPU registers. You can run it at any breakpoint.
- `bt` shows you a backtrace of the stack, which includes the nested functions you have executed. For example, if you executed `A()` and inside `A()` you executed `B()`, it will show you `B → A`.

Suggested reading

RISC-V GDB Tutorial, SHAKTI Development Team

4 Miscellaneous

4.1 Understanding the xv6 Code Directory Structure

The xv6 OS from the GitHub repository has the following 3 directories in the root directory.

1. `kernel/` contains assembly and C code files that belong to the kernel. Some important and noteworthy files include the following two files: (1) `start.S` which contains first assembly code that runs when xv6 boots, and (2) `main.c` which contains first C code that runs when xv6 boots.
2. `user/` contains code files related to user programs (e.g., `ls`, `cat`, etc.). Check out the code for `'ls.c'` if you're interested.
3. `mkfs/` contains code files for building an xv6 file system and image.

Suggested reading

RISC-V Assembly Language, Stephen Marz